



Santosh Sam Koshy, JD C-DAC, Hyderabad santoshsk@cdac.in

AGENDA



- Recapitulation of previous lecture
- Learning Outcomes for this topic
- Introduction
- ioctl implementation

RECAPITULATION



- In previous lecture, we learnt how to build a simple char device driver
- How to read from and write to the device
- Manipulating different kernel data structures

INTRODUCTION



- Only reading to and writing from the device is not enough
- Drivers need some hardware control mechnism also
- The *ioctl* method provides a mechanism to perform control actions on the device through the user space as well as in kernel space

Prototype of *ioctl*



• In user space, *ioctl* system call has following prototype:

int ioctl (int fd, unsigned long cmd, ...)

- fd is file descriptor of the device file
- cmd is used to select the behavior of the device
- Third argument is an optional argument with no type check during compilation
- In kernel space, prototype for *ioctl* is:

long (*ioctl) (struct file *filp, unsigned int cmd, unsigned long arp)

- inode and filp pointers are the values corresponding to the file descriptor fd
- cmd argument is passed from the user unchanged
- Optional argument is passed in the form of an unsigned long
- If invoking program is not passing third argument, the arg value received will be undefined
- In the file_operations structure, specify 'unlocked_ioctl' to implement functionality

ioctl commands



- *ioctl* commands are represented using some numbers
- These numbers should be unique to prevent an unwanted operation on a device mistakenly:
 - To find a unique number see Documentation/userspace-api/ioctl/ioctl-number.rst
 - This text lists the magic numbers which is used throughout the kernel
- This number is 32 bits wide which includes 4 bitfields:
 - **type(magic number):** choose one number only and use through out your program. 8 bits
 - Ordinal number: sequential number, 8bits wide
 - **Direction:** direction of data transfer, if particular command involves data transfer. 2 bits wide
 - The possible values are: _IOC_NONE, _IOC_READ, _IOC_WRITE AND _IOC_READ|_IOC_WRITE
 - Size: the size of user data involved. It's 14 bits wide
 - These are shared between the user application and the kernel driver.

Magic Numbers -Internally



Bits	Meaning
31-30:	
00	no parameters: uses _IO macro
10	read: _IOR
01	write: _IOW
11	read/write: _IOWR
29-16	size of arguments
15-8	ascii character supposedly unique to each driver
7-0	Ordinal Number/Unique Command Differentiation

Creating *ioctl* Commands



- IOCTL commands can be defined in a header file using special functions that aggregate the bit fields into unique commands
 - _IO(type, nr): command that has no argument
 - _IOR(type, nr, datatype): for reading data from driver
 - _IOW(type, nr, datatype): for writing data
 - _IOWR(type, nr, datatype: for bidirectional transfers
- This header file is required by the application and the kernel driver that uses the command
- Types and ordinal number is passed as argument and size of the datatype can be passed directly
 - DO NOT PASS sizeof(type) the 3rd argument is already replaced by sizeof() type

```
/* Using 'k' as magic number */
/* Refer Documentation/ioctl/ioctl-numbers.txt for free magic number */
#define serial_magic 'k'

/* Defining the ioctl commands */
#define serial_reset _IO(serial_magic, 0)
#define serial_baud_set _IOW(serial_magic, 1, int)
#define serial_status _IOR(serial_magic, 2, char)
#define serial_control _IOWR(serial_magic, 3, char)
#define serial_maxnr 4
~
```

ioctl – Kernel Level Implementation



```
long (*ioctl) (struct file *filp, unsigned int cmd, unsigned long arp)
{ //Implementation of ioctl at kernel level is usually a switch statement based on command number
     Switch (cmd) {
          Case BAUDRATE:
                //Do something here to set the baud rate....
                Break:
          Case PARITY:
               //Do something here to set the parity ....
                Break;
          Case STOP BITS:
                //Do something here to set the stop bits....
                Break;
          Case All:
               //Do something here to set the all the data using a user defined structure.....
                Break:
          Default: The default returns -EINVAL("Invalid argument")
               //Error – No such command
               Break;
```

Implementation of *ioctl* Commands at kernel level



• Simple switch case statements are used to implement at kernel level

```
/* for storing the baud value */
unsigned int BAUD VALUE;
unsigned char STATUS, CONTROL; /* used as a temporary buffer */
int retval:
                              /* for storing return value */
/* arg is the data sent from the user space */
switch(cmd)
       case serial baud: // setting the baud rate of serial device
               retval = get user(BAUD VALUE, (int user *)arg); /* *arg -> BAUD VALUE */
               outw(BAUD VALUE, address baud reg); /* baud value is written to device baud register */
               break:
       case status:
                      // getting the status of serial device
                                                  /* status is read from the device status register */
               STATUS = inb(address status reg);
               retval = put user(STATUS, (char user *)arg); /* STATUS -> *arg */
               break:
       case serial control:
                              // getting the status result
               retval = get user(CONTROL, (char user *)arg); /* *arg -> CONTROL*/
               outb(CONTROL, address control reg); /* control bit writing to the device */
                                      /* if writing to the device is successful */
               if(retval == 0)
                       CONTROL = inb(address control reg); /* returning the current control value */
                       retval = put user(CONTROL, (char user *)arg); /* CONTROL -> *arg */
               break:
```

Optimised Data Transfer



- Since in *ioctl*, data transfer is rarely more than 8 bytes; in such case following functions works well:
- put_user(data, ptr) and __put_user(data, ptr): this macro write data to user space
 - Relatively faster than *copy_to_user()*
 - Size of data transfer depends on type of *ptr* argument
 - __put_user() performs less checking
- Similarly, get_user(data, ptr) and __get_user(data, ptr)

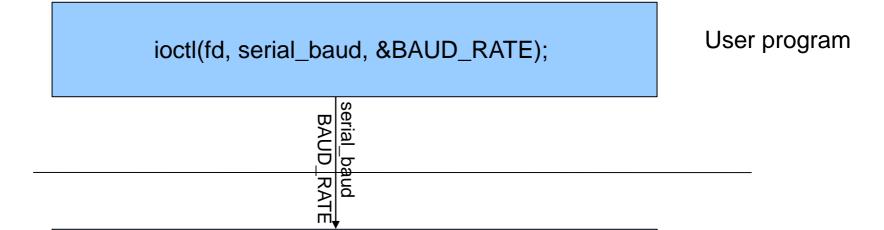
Implementation of *ioctl* at user-space



• It's just calling the function with appropriate cmd arguments

ioctl quick view





case serial_baud: // setting the baud rate of serial device
retval = __get_user(BAUD_VALUE, (int __user *)arg);
outw(BAUD_VALUE, address_baud_reg);

Driver

baud_reg = BAUD_VALUE

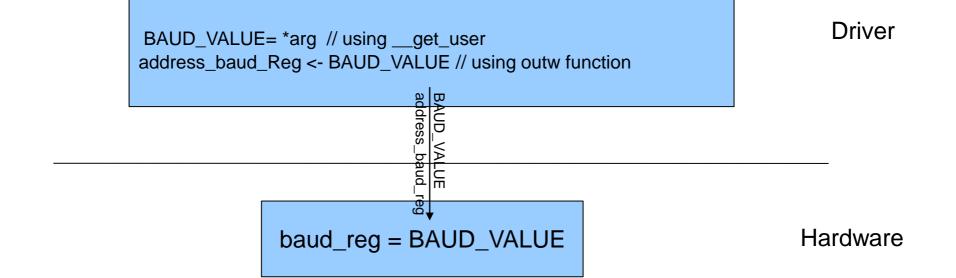
Hardware

ioctl quick view



ioctl(fd, serial_baud, &BAUD_RATE);

User program





Thank you

Centre for Development of Advanced Computing, Hyderabad



Centre for Development of Advanced Computing, Hyderabad